

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import math
from pandas.api.types import CategoricalDtype
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, roc_curve, roc_auc_score
from sklearn.model_selection import GridSearchCV, KFold
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import BaggingClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import StandardScaler
from statsmodels.stats.outliers_influence import variance_inflation_factor
```

```
In [2]: train_set = pd.read_csv('http://archive.ics.uci.edu/ml/machine-learning-databases/adult/adult.data', header = None)
test_set = pd.read_csv('http://archive.ics.uci.edu/ml/machine-learning-databases/adult/adult.test', skiprows = 1, header = None)
col_labels = ['age', 'workclass', 'fnlwgt', 'education', 'education_num', 'marital_status', 'occupation', 'relationship',
              'race', 'sex', 'capital_gain', 'capital_loss', 'hours_per_week', 'native_country', 'wage_class']
train_set.columns = col_labels
test_set.columns = col_labels
```

```
In [3]: print('Train size:: {}'.format(train_set.shape))
print('Test size:: {}'.format(test_set.shape))
```

```
Train size:: (32561, 15)
Test size:: (16281, 15)
```

```
In [4]: train_set.head()
```

Out[4]:

	age	workclass	fnlwgt	education	education_num	marital_status	occupation	relationship
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife

In [5]: `test_set.head()`

Out[5]:

	age	workclass	fnlwgt	education	education_num	marital_status	occupation	relationship
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-child
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband
3	44	Private	160323	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband
4	18	?	103497	Some-college	10	Never-married	?	Own-child

In [6]: `train_set.describe()`

Out[6]:

	age	fnlwgt	education_num	capital_gain	capital_loss	hours_per_week
count	32561.000000	3.256100e+04	32561.000000	32561.000000	32561.000000	32561.000000
mean	38.581647	1.897784e+05	10.080679	1077.648844	87.303830	40.437456
std	13.640433	1.055500e+05	2.572720	7385.292085	402.960219	12.347429
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	1.000000
25%	28.000000	1.178270e+05	9.000000	0.000000	0.000000	40.000000
50%	37.000000	1.783560e+05	10.000000	0.000000	0.000000	40.000000
75%	48.000000	2.370510e+05	12.000000	0.000000	0.000000	45.000000
max	90.000000	1.484705e+06	16.000000	99999.000000	4356.000000	99.000000

In [7]: `train_set.isna().sum()`

Out[7]:

```
age          0
workclass    0
fnlwgt       0
education    0
education_num 0
marital_status 0
occupation   0
relationship 0
race         0
sex          0
capital_gain 0
capital_loss 0
hours_per_week 0
native_country 0
wage_class   0
dtype: int64
```

```
In [8]: train_set.describe(include=["object"]).T
```

Out[8]:

	count	unique	top	freq
workclass	32561	9	Private	22696
education	32561	16	HS-grad	10501
marital_status	32561	7	Married-civ-spouse	14976
occupation	32561	15	Prof-specialty	4140
relationship	32561	6	Husband	13193
race	32561	5	White	27816
sex	32561	2	Male	21790
native_country	32561	42	United-States	29170
wage_class	32561	2	<=50K	24720

```
In [9]: test_set.describe(include=["object"]).T
```

Out[9]:

	count	unique	top	freq
workclass	16281	9	Private	11210
education	16281	16	HS-grad	5283
marital_status	16281	7	Married-civ-spouse	7403
occupation	16281	15	Prof-specialty	2032
relationship	16281	6	Husband	6523
race	16281	5	White	13946
sex	16281	2	Male	10860
native_country	16281	41	United-States	14662
wage_class	16281	2	<=50K	12435

```
In [10]: # Check if duplicate data exist
check_dup = train_set.duplicated().any()
print("Are there any duplicated values in data? ", check_dup)

if check_dup:
    train_set = train_set.drop_duplicates()
else:
    print("There are not duplicated values in data.")

check_dup_test = test_set.duplicated().any()
print("Are there any duplicated values in data? ", check_dup)

if check_dup:
    test_set = test_set.drop_duplicates()
else:
    print("There are not duplicated values in data.")
```

Are there any duplicated values in data? True

Are there any duplicated values in data? True

```
In [11]: object_columns=train_set.select_dtypes(include=["object"]).columns
         for i in range(len(object_columns)):
             print("----- {}-----".format(object_columns[i]))
             print(train_set[object_columns[i]].value_counts())
```

----- workclass-----

Private	22673
Self-emp-not-inc	2540
Local-gov	2093
?	1836
State-gov	1298
Self-emp-inc	1116
Federal-gov	960
Without-pay	14
Never-worked	7

Name: workclass, dtype: int64

----- education-----

HS-grad	10494
Some-college	7282
Bachelors	5353
Masters	1722
Assoc-voc	1382
11th	1175
Assoc-acdm	1067
10th	933
7th-8th	645
Prof-school	576
9th	514
12th	433
Doctorate	413
5th-6th	332
1st-4th	166
Preschool	50

Name: education, dtype: int64

----- marital_status-----

Married-civ-spouse	14970
Never-married	10667
Divorced	4441
Separated	1025
Widowed	993
Married-spouse-absent	418
Married-AF-spouse	23

Name: marital_status, dtype: int64

----- occupation-----

Prof-specialty	4136
Craft-repair	4094
Exec-managerial	4065
Adm-clerical	3768
Sales	3650
Other-service	3291
Machine-op-inspct	2000
?	1843
Transport-moving	1597
Handlers-cleaners	1369
Farming-fishing	992
Tech-support	927
Protective-serv	649
Priv-house-serv	147
Armed-Forces	9

Name: occupation, dtype: int64

----- relationship-----

Husband	13187
---------	-------

```

Not-in-family      8292
Own-child          5064
Unmarried          3445
Wife               1568
Other-relative     981
Name: relationship, dtype: int64
----- race-----
White              27795
Black              3122
Asian-Pac-Islander 1038
Amer-Indian-Eskimo 311
Other              271
Name: race, dtype: int64
----- sex-----
Male              21775
Female           10762
Name: sex, dtype: int64
----- native_country-----
United-States      29153
Mexico             639
?                  582
Philippines        198
Germany            137
Canada            121
Puerto-Rico       114
El-Salvador        106
India              100
Cuba               95
England            90
Jamaica            81
South              80
China              75
Italy              73
Dominican-Republic 70
Vietnam            67
Guatemala          62
Japan              62
Poland             60
Columbia           59
Taiwan             51
Haiti              44
Iran               43
Portugal           37
Nicaragua          34
Peru               31
France             29
Greece             29
Ecuador            28
Ireland            24
Hong               20
Trinidad&Tobago    19
Cambodia           19
Thailand           18
Laos               18
Yugoslavia         16
Outlying-US(Guam-USVI-etc) 14
Hungary            13

```

```

Honduras      13
Scotland      12
Holand-Netherlands  1
Name: native_country, dtype: int64
----- wage_class-----
<=50K      24698
>50K       7839
Name: wage_class, dtype: int64

```

```

In [12]: train_set.replace([' ?'], np.nan, inplace=True)
         test_set.replace([' ?'], np.nan, inplace=True)

```

```

In [13]: train_set.isnull().sum()

```

```

Out[13]: age      0
         workclass  1836
         fnlwgt    0
         education  0
         education_num  0
         marital_status  0
         occupation  1843
         relationship  0
         race      0
         sex      0
         capital_gain  0
         capital_loss  0
         hours_per_week  0
         native_country  582
         wage_class    0
         dtype: int64

```

```

In [14]: test_set.isnull().sum()

```

```

Out[14]: age      0
         workclass  963
         fnlwgt    0
         education  0
         education_num  0
         marital_status  0
         occupation  966
         relationship  0
         race      0
         sex      0
         capital_gain  0
         capital_loss  0
         hours_per_week  0
         native_country  274
         wage_class    0
         dtype: int64

```



```
In [15]: object_columns=train_set.select_dtypes(include=["object"]).columns
         for i in range(len(object_columns)):
             print("----- {}-----".format(object_columns[i]))
             print(train_set[object_columns[i]].value_counts())
```

```

----- workclass-----
Private                22673
Self-emp-not-inc      2540
Local-gov             2093
State-gov             1298
Self-emp-inc          1116
Federal-gov           960
Without-pay           14
Never-worked           7
Name: workclass, dtype: int64
----- education-----
HS-grad               10494
Some-college          7282
Bachelors             5353
Masters               1722
Assoc-voc             1382
11th                  1175
Assoc-acdm            1067
10th                  933
7th-8th               645
Prof-school           576
9th                   514
12th                  433
Doctorate             413
5th-6th               332
1st-4th               166
Preschool             50
Name: education, dtype: int64
----- marital_status-----
Married-civ-spouse    14970
Never-married         10667
Divorced              4441
Separated             1025
Widowed               993
Married-spouse-absent  418
Married-AF-spouse      23
Name: marital_status, dtype: int64
----- occupation-----
Prof-specialty        4136
Craft-repair          4094
Exec-managerial       4065
Adm-clerical          3768
Sales                 3650
Other-service         3291
Machine-op-inspct     2000
Transport-moving      1597
Handlers-cleaners     1369
Farming-fishing       992
Tech-support          927
Protective-serv       649
Priv-house-serv       147
Armed-Forces          9
Name: occupation, dtype: int64
----- relationship-----
Husband               13187
Not-in-family         8292
Own-child             5064

```

```

Unmarried          3445
Wife                1568
Other-relative     981
Name: relationship, dtype: int64
----- race-----
White              27795
Black              3122
Asian-Pac-Islander 1038
Amer-Indian-Eskimo 311
Other              271
Name: race, dtype: int64
----- sex-----
Male              21775
Female           10762
Name: sex, dtype: int64
----- native_country-----
United-States      29153
Mexico              639
Philippines        198
Germany            137
Canada             121
Puerto-Rico        114
El-Salvador         106
India              100
Cuba                95
England             90
Jamaica             81
South              80
China              75
Italy              73
Dominican-Republic 70
Vietnam            67
Japan              62
Guatemala          62
Poland             60
Columbia           59
Taiwan             51
Haiti              44
Iran               43
Portugal           37
Nicaragua          34
Peru               31
Greece             29
France             29
Ecuador            28
Ireland            24
Hong               20
Cambodia           19
Trinidad&Tobago    19
Thailand           18
Laos               18
Yugoslavia         16
Outlying-US(Guam-USVI-etc) 14
Hungary            13
Honduras           13
Scotland           12
Holand-Netherlands 1

```

Name: native_country, dtype: int64

----- wage_class-----

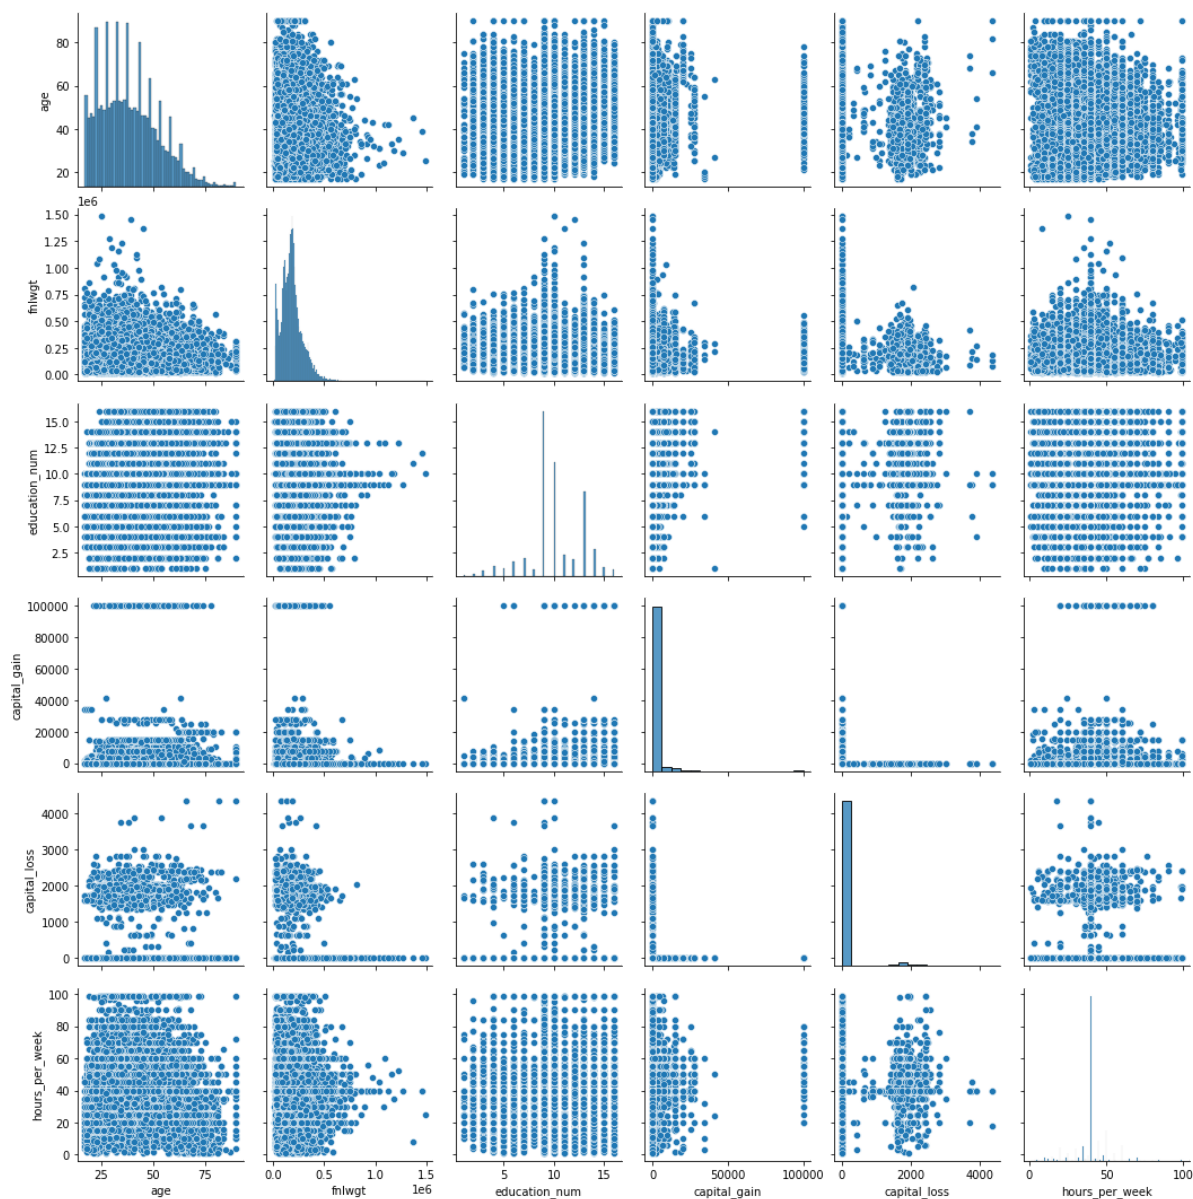
<=50K 24698

>50K 7839

Name: wage_class, dtype: int64

In [16]: sns.pairplot(train_set)

Out[16]: <seaborn.axisgrid.PairGrid at 0x214feef3048>



```
In [17]: train_set.education.value_counts()
```

```
Out[17]: HS-grad      10494
Some-college  7282
Bachelors    5353
Masters      1722
Assoc-voc    1382
11th         1175
Assoc-acdm   1067
10th         933
7th-8th      645
Prof-school  576
9th          514
12th         433
Doctorate    413
5th-6th      332
1st-4th      166
Preschool    50
Name: education, dtype: int64
```

```
In [18]: train_set.education_num.value_counts()
```

```
Out[18]: 9      10494
10      7282
13      5353
14      1722
11      1382
7       1175
12      1067
6       933
4       645
15      576
5       514
8       433
16      413
3       332
2       166
1        50
Name: education_num, dtype: int64
```

```
In [19]: ## As education and education_num are same, we are dropping education column
train_set.drop(columns=['education'], inplace=True)
test_set.drop(columns=['education'], inplace=True)
```

```
In [20]: train_set["education_num"]=train_set["education_num"].astype(CategoricalDtype(
ordered=True)) #ordinal data.
test_set["education_num"]=test_set["education_num"].astype(CategoricalDtype(ordered=True))
train_set["education_num"].head()
```

```
Out[20]: 0    13
1    13
2     9
3     7
4    13
Name: education_num, dtype: category
Categories (16, int64): [1 < 2 < 3 < 4 ... 13 < 14 < 15 < 16]
```

```
In [21]: train_set.education_num.value_counts()
```

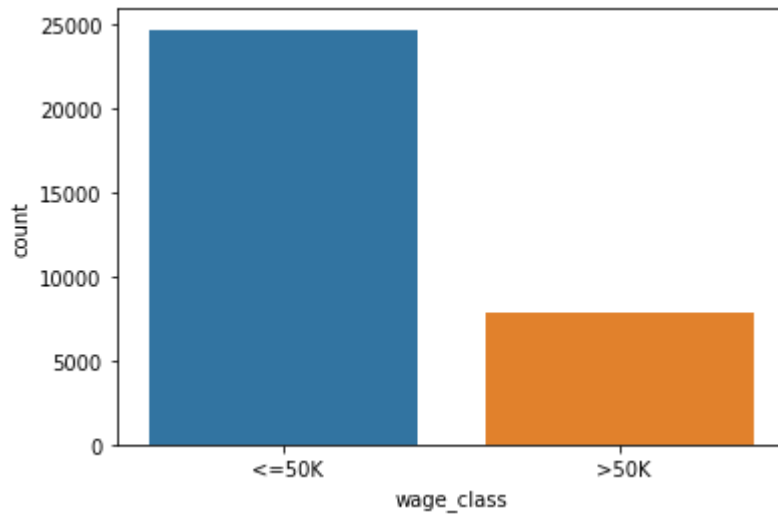
```
Out[21]: 9    10494
10    7282
13    5353
14    1722
11    1382
7     1175
12    1067
6      933
4      645
15     576
5      514
8      433
16     413
3      332
2      166
1        50
Name: education_num, dtype: int64
```

```
In [22]: sns.countplot(train_set.wage_class)
```

C:\Users\Urvi\AppData\Roaming\Python\Python37\site-packages\seaborn_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

```
Out[22]: <AxesSubplot:xlabel='wage_class', ylabel='count'>
```



```
In [23]: train_set.skew()
```

```
Out[23]: age          0.557663  
fnlwgt        1.447703  
capital_gain  11.949403  
capital_loss   4.592702  
hours_per_week 0.228759  
dtype: float64
```

```
In [24]: plt.figure(figsize=(30,15))

plt.subplot(321)
sns.boxplot(train_set["age"])
plt.xticks(rotation=90)

plt.subplot(322)
sns.boxplot(train_set["capital_loss"])
plt.xticks(rotation=90)

plt.subplot(323)
sns.boxplot(train_set["capital_gain"])
plt.xticks(rotation=90)

plt.subplot(324)
sns.boxplot(train_set["fnlwt"])
plt.xticks(rotation=90)

plt.subplot(325)
sns.boxplot(train_set["hours_per_week"])
plt.xticks(rotation=90)

plt.subplots_adjust(hspace=0.5)
plt.show()
```


C:\Users\Urvi\AppData\Roaming\Python\Python37\site-packages\seaborn_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

C:\Users\Urvi\AppData\Roaming\Python\Python37\site-packages\seaborn_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

C:\Users\Urvi\AppData\Roaming\Python\Python37\site-packages\seaborn_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

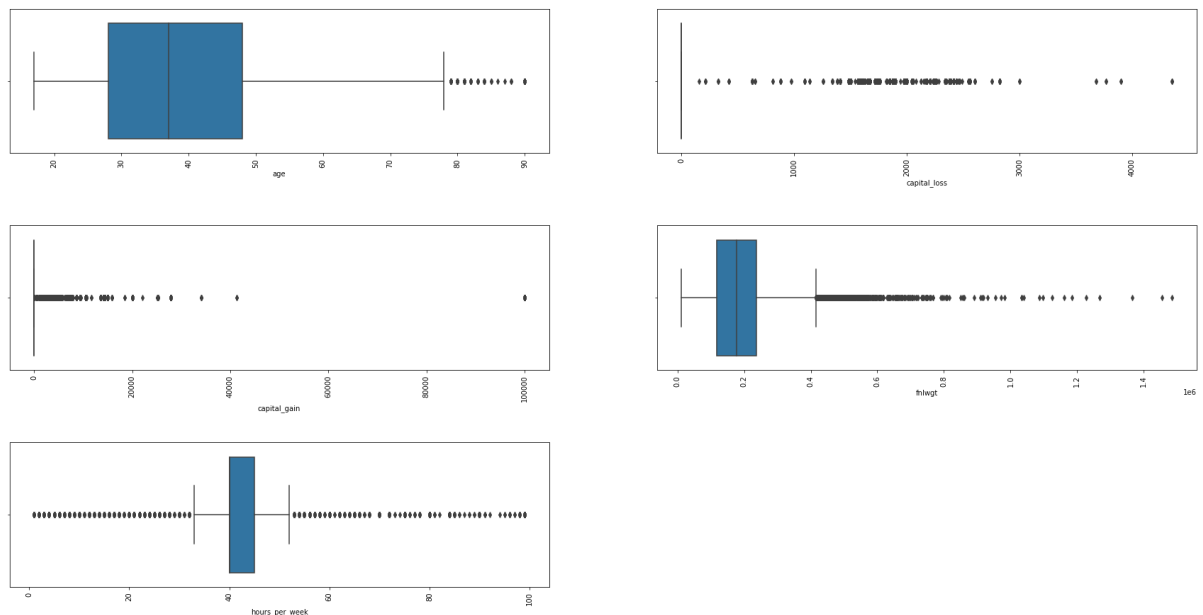
FutureWarning

C:\Users\Urvi\AppData\Roaming\Python\Python37\site-packages\seaborn_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

C:\Users\Urvi\AppData\Roaming\Python\Python37\site-packages\seaborn_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning




```
In [25]: numeric_columns=list(train_set.select_dtypes(include=["int64"]).columns) #numeric columns
numeric_columns
```

```
Out[25]: ['age', 'fnlwgt', 'capital_gain', 'capital_loss', 'hours_per_week']
```


```
In [26]: train_set["occupation"].value_counts().plot.barh(color="red")
```

Out[26]: <AxesSubplot:>



```
In [27]: train_set["workclass"].value_counts().plot.barh(color="orange")
```

Out[27]: <AxesSubplot:>



```
In [28]: train_set.native_country.value_counts()
```

```
Out[28]: United-States      29153
Mexico      639
Philippines 198
Germany     137
Canada      121
Puerto-Rico 114
El-Salvador 106
India       100
Cuba        95
England     90
Jamaica     81
South       80
China       75
Italy       73
Dominican-Republic 70
Vietnam     67
Japan       62
Guatemala   62
Poland      60
Columbia    59
Taiwan      51
Haiti       44
Iran        43
Portugal    37
Nicaragua   34
Peru        31
Greece      29
France      29
Ecuador     28
Ireland     24
Hong        20
Cambodia    19
Trinidad&Tobago 19
Thailand     18
Laos        18
Yugoslavia  16
Outlying-US(Guam-USVI-etc) 14
Hungary     13
Honduras    13
Scotland    12
Holand-Netherlands 1
Name: native_country, dtype: int64
```

```
In [29]: for i in ["occupation", "workclass", "native_country"]:
          train_set[i].fillna(train_set[i].mode()[0], inplace=True)
          test_set[i].fillna(test_set[i].mode()[0], inplace=True)
```

```
In [30]: train_set.isnull().sum()
```

```
Out[30]: age                0  
workclass            0  
fnlwgt              0  
education_num       0  
marital_status      0  
occupation          0  
relationship        0  
race                0  
sex                 0  
capital_gain        0  
capital_loss        0  
hours_per_week      0  
native_country      0  
wage_class          0  
dtype: int64
```

```
In [31]: test_set.isnull().sum()
```

```
Out[31]: age                0  
workclass            0  
fnlwgt              0  
education_num       0  
marital_status      0  
occupation          0  
relationship        0  
race                0  
sex                 0  
capital_gain        0  
capital_loss        0  
hours_per_week      0  
native_country      0  
wage_class          0  
dtype: int64
```

```
In [32]: #Assigning the numeric values to the string type variables
number = LabelEncoder()
train_set['workclass'] = number.fit_transform(train_set['workclass'])
train_set['marital_status'] = number.fit_transform(train_set['marital_status'])
train_set['occupation'] = number.fit_transform(train_set['occupation'])
train_set['relationship'] = number.fit_transform(train_set['relationship'])
train_set['race'] = number.fit_transform(train_set['race'])
train_set['sex'] = number.fit_transform(train_set['sex'])
train_set['native_country'] = number.fit_transform(train_set['native_country'])
train_set['wage_class'] = number.fit_transform(train_set['wage_class'])

test_set['workclass'] = number.fit_transform(test_set['workclass'])
test_set['marital_status'] = number.fit_transform(test_set['marital_status'])
test_set['occupation'] = number.fit_transform(test_set['occupation'])
test_set['relationship'] = number.fit_transform(test_set['relationship'])
test_set['race'] = number.fit_transform(test_set['race'])
test_set['sex'] = number.fit_transform(test_set['sex'])
test_set['native_country'] = number.fit_transform(test_set['native_country'])
test_set['wage_class'] = number.fit_transform(test_set['wage_class'])
```

```
In [33]: train_set['age_bin'] = pd.cut(train_set['age'], 20)
test_set['age_bin'] = pd.cut(test_set['age'], 20)
```

```
In [34]: train_set['hours-per-week_bin'] = pd.cut(train_set['hours_per_week'], 10)
test_set['hours-per-week_bin'] = pd.cut(test_set['hours_per_week'], 10)
```

```
In [35]: train_set[['wage_class', 'age']].groupby(['wage_class'], as_index=False).mean()
.sort_values(by='age', ascending=False)
```

Out[35]:

	wage_class	age
1	1	44.250925
0	0	36.787392

```
In [36]: train_set = train_set.apply(LabelEncoder().fit_transform)
train_set.head()
```

Out[36]:

	age	workclass	fnlwgt	education_num	marital_status	occupation	relationship	race	sex	c
0	22	6	2671	12	4	0	1	4	1	
1	33	5	2926	12	2	3	0	4	1	
2	21	3	14086	8	0	5	1	4	1	
3	36	3	15336	6	2	5	0	2	1	
4	11	3	19355	12	2	9	5	2	0	

```
In [37]: test_set = test_set.apply(LabelEncoder().fit_transform)
test_set.head()
```

Out[37]:

	age	workclass	fnlwgt	education_num	marital_status	occupation	relationship	race	sex	c
0	8	3	8931	6	4	6	3	2	1	
1	21	3	1888	8	2	4	0	4	1	
2	11	1	11540	11	2	10	0	4	1	
3	27	3	5146	9	2	6	0	2	1	
4	1	3	2450	9	4	9	3	4	0	

```
In [38]: train_x = train_set.drop(columns=['wage_class', 'age', 'hours_per_week'])
test_x = test_set.drop(columns=['wage_class', 'age', 'hours_per_week'])
train_y = train_set.wage_class
test_y = test_set.wage_class
```

Checking VIF

```
In [39]: scalar = StandardScaler()
X_scaled = scalar.fit_transform(train_x)
```

```
In [40]: vif = pd.DataFrame()
vif["vif"] = [variance_inflation_factor(X_scaled,i) for i in range(X_scaled.shape[1])]
vif["Features"] = train_x.columns

#Let's check the values
vif
```

Out[40]:

	vif	Features
0	1.010422	workclass
1	1.014561	fnlwgt
2	1.076122	education_num
3	1.123101	marital_status
4	1.009878	occupation
5	1.678307	relationship
6	1.032643	race
7	1.558542	sex
8	1.056592	capital_gain
9	1.022096	capital_loss
10	1.025766	native_country
11	1.165779	age_bin
12	1.118631	hours-per-week_bin

XGBoost Classifier

```
In [41]: # fit model no training data
model = XGBClassifier(objective='binary:logistic')
model.fit(train_x, train_y)
```

```
Out[41]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
colsample_bynode=1, colsample_bytree=1, gamma=0, gpu_id=-1,
importance_type='gain', interaction_constraints='',
learning_rate=0.300000012, max_delta_step=0, max_depth=6,
min_child_weight=1, missing=nan, monotone_constraints='()',
n_estimators=100, n_jobs=0, num_parallel_tree=1,
objective='binary:logistic', random_state=0, reg_alpha=0,
reg_lambda=1, scale_pos_weight=1, subsample=1,
tree_method='exact', validate_parameters=1, verbosity=None)
```

```
In [42]: # cheking training accuracy
y_pred = model.predict(train_x)
predictions = [round(value) for value in y_pred]
accuracy = accuracy_score(train_y,predictions)
accuracy
```

Out[42]: 0.9019270369118234

```
In [43]: # cheking initial test accuracy
y_pred = model.predict(test_x)
predictions = [round(value) for value in y_pred]
accuracy = accuracy_score(test_y,predictions)
accuracy
```

Out[43]: 0.854939788645859

```
In [44]: param_grid={
    'learning_rate': [1, 0.5, 0.1, 0.01, 0.001],
    'max_depth': [3, 5, 10, 20],
    'n_estimators': [10, 50, 100, 200]
}
```

```
In [45]: grid= GridSearchCV(XGBClassifier(objective='binary:logistic'), param_grid)
```

```
In [46]: grid.fit(train_x, train_y)
```

```
Out[46]: GridSearchCV(cv=None, error_score=nan,
                      estimator=XGBClassifier(base_score=None, booster=None,
                                              colsample_bylevel=None,
                                              colsample_bynode=None,
                                              colsample_bytree=None, gamma=None,
                                              gpu_id=None, importance_type='gain',
                                              interaction_constraints=None,
                                              learning_rate=None, max_delta_step=None,
                                              max_depth=None, min_child_weight=None,
                                              missing=nan, monotone_constraints=None,
                                              n_estimators=200, n_jobs=1,
                                              random_state=None, reg_alpha=None,
                                              reg_lambda=None, scale_pos_weight=None,
                                              subsample=None, tree_method=None,
                                              validate_parameters=None, verbosity=None),
                      iid='deprecated', n_jobs=None,
                      param_grid={'learning_rate': [1, 0.5, 0.1, 0.01, 0.001],
                                  'max_depth': [3, 5, 10, 20],
                                  'n_estimators': [10, 50, 100, 200]},
                      pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                      scoring=None, verbose=0)
```

```
In [47]: grid.best_params_
```

Out[47]: {'learning_rate': 0.1, 'max_depth': 5, 'n_estimators': 200}


```
In [48]: # Create new model using the same parameters
new_model=XGBClassifier(learning_rate= 0.1, max_depth= 5, n_estimators= 200)
new_model.fit(train_x, train_y)
```

```
Out[48]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                      colsample_bynode=1, colsample_bytree=1, gamma=0, gpu_id=-1,
                      importance_type='gain', interaction_constraints='',
                      learning_rate=0.1, max_delta_step=0, max_depth=5,
                      min_child_weight=1, missing=nan, monotone_constraints='()',
                      n_estimators=200, n_jobs=0, num_parallel_tree=1,
                      objective='binary:logistic', random_state=0, reg_alpha=0,
                      reg_lambda=1, scale_pos_weight=1, subsample=1,
                      tree_method='exact', validate_parameters=1, verbosity=None)
```

```
In [49]: y_pred_new = new_model.predict(test_x)
predictions_new = [round(value) for value in y_pred_new]
accuracy_new_xgboost = accuracy_score(test_y, predictions_new)
accuracy_new_xgboost
```

```
Out[49]: 0.8597321209142296
```

```
In [50]: new_model.get_booster().get_score(importance_type="gain")
```

```
Out[50]: {'relationship': 128.6458774135325,
          'education_num': 41.5177175380451,
          'capital_gain': 37.861184960465046,
          'capital_loss': 18.21865458632269,
          'age_bin': 15.75366145925717,
          'occupation': 9.55463262210465,
          'hours-per-week_bin': 12.643608384179403,
          'fnlwgt': 3.129041437650815,
          'marital_status': 38.846798557440955,
          'sex': 9.694870202,
          'race': 4.3695650333863645,
          'workclass': 5.375567432009281,
          'native_country': 3.365511446216217}
```


Features relationship, education_num, marital_status, capital_gain, capital_loss, age_bin, sex, hours-per-week_bin has more importance as compared to other variables

```
In [51]: y_pred_xgboost = new_model.predict(test_x)
```

```
In [52]: conf_mat_xgboost = confusion_matrix(test_y, y_pred_xgboost)
true_positive_xgboost = conf_mat_xgboost[0][0]
false_positive_xgboost = conf_mat_xgboost[0][1]
false_negative_xgboost = conf_mat_xgboost[1][0]
true_negative_xgboost = conf_mat_xgboost[1][1]
Accuracy_xgboost = (true_positive_xgboost + true_negative_xgboost) / (true_pos
itive_xgboost + false_positive_xgboost + false_negative_xgboost + true_negative
_xgboost)
Precision_xgboost = true_positive_xgboost / (true_positive_xgboost + false_positiv
e_xgboost)
Recall_xgboost = true_positive_xgboost / (true_positive_xgboost + false_negative_x
gboost)
F1_Score_xgboost = 2 * (Recall_xgboost * Precision_xgboost) / (Recall_xgboost +
Precision_xgboost)
auc_xgboost = roc_auc_score(test_y, y_pred_xgboost)
fpr_xgboost, tpr_xgboost, thresholds_xgboost = roc_curve(test_y, y_pred_xgboos
t)
print('Accuracy:: ', Accuracy_xgboost)
print('Precision:: ', Precision_xgboost)
print('Recall:: ', Recall_xgboost)
print('F1 Score:: ', F1_Score_xgboost)
print('AUC:: ', auc_xgboost)
```

```
Accuracy:: 0.8597321209142296
Precision:: 0.9432823813354787
Recall:: 0.8813801398180862
F1 Score:: 0.9112812342128784
AUC:: 0.7664929847395022
```

```
In [53]: plt.plot(fpr_xgboost, tpr_xgboost, color='orange', label='ROC')
plt.plot([0, 1], [0, 1], color='darkblue', linestyle='--', label='ROC curve (a
rea = %0.2f)' % auc_xgboost)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend()
plt.show()
```



Logistic Regression

```
In [54]: log_reg = LogisticRegression(max_iter=300)
log_reg.fit(train_x, train_y)
```

```
Out[54]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, l1_ratio=None, max_iter=300,
multi_class='auto', n_jobs=None, penalty='l2',
random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
warm_start=False)
```

```
In [55]: accuracy_score(train_y, log_reg.predict(train_x))
```

```
Out[55]: 0.8101238589913022
```

```
In [56]: y_pred_log = log_reg.predict(test_x)
```


```
In [57]: accuracy_score(test_y, y_pred_log)
```

```
Out[57]: 0.8109486360285082
```

```
In [58]: conf_mat = confusion_matrix(test_y, y_pred_log)
true_positive = conf_mat[0][0]
false_positive = conf_mat[0][1]
false_negative = conf_mat[1][0]
true_negative = conf_mat[1][1]
Accuracy = (true_positive + true_negative) / (true_positive + false_positive +
false_negative + true_negative)
Precision = true_positive / (true_positive + false_positive)
Recall = true_positive / (true_positive + false_negative)
F1_Score = 2 * (Recall * Precision) / (Recall + Precision)
auc = roc_auc_score(test_y, y_pred_log)
fpr_log, tpr_log, thresholds_log = roc_curve(test_y, y_pred_log)
print('Accuracy:: ', Accuracy)
print('Precision:: ', Precision)
print('Recall:: ', Recall)
print('F1 Score:: ', F1_Score)
print('AUC:: ', auc)
```

```
Accuracy:: 0.8109486360285082
Precision:: 0.9320193081255028
Recall:: 0.838459868278208
F1 Score:: 0.882767554387168
AUC:: 0.6758380471984768
```

```
In [59]: plt.plot(fpr_log, tpr_log, color='orange', label='ROC')
plt.plot([0, 1], [0, 1], color='darkblue', linestyle='--', label='ROC curve (a
rea = %0.2f)' % auc)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend()
plt.show()
```



Decision Tree Classifier

```
In [60]: clf = DecisionTreeClassifier()
clf.fit(train_x, train_y)
```

```
Out[60]: DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',
                                max_depth=None, max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, presort='deprecated',
                                random_state=None, splitter='best')
```

```
In [61]: clf.score(train_x, train_y)
```

```
Out[61]: 0.9998770630359284
```

```
In [62]: y_pred_dtc = clf.predict(test_x)
```

```
In [63]: clf.score(test_x, test_y)
```

```
Out[63]: 0.8078151880068813
```

```
In [64]: grid_param_dt = {
    'criterion': ['gini', 'entropy'],
    'max_depth' : range(2,32,1),
    'min_samples_leaf' : range(1,10,1),
    'min_samples_split': range(2,10,1),
    'splitter' : ['best', 'random']
}
```

```
In [65]: grid_search_decision = GridSearchCV(estimator = clf, param_grid = grid_param_d
t, cv = 5, n_jobs = -1)
```

```
In [66]: grid_search_decision.fit(train_x, train_y)
```

```
Out[66]: GridSearchCV(cv=5, error_score=nan,
                      estimator=DecisionTreeClassifier(ccp_alpha=0.0, class_weight=Non
e,
                                                         criterion='gini', max_depth=Non
e,
                                                         max_features=None,
                                                         max_leaf_nodes=None,
                                                         min_impurity_decrease=0.0,
                                                         min_impurity_split=None,
                                                         min_samples_leaf=1,
                                                         min_samples_split=2,
                                                         min_weight_fraction_leaf=0.0,
                                                         presort='deprecated',
                                                         random_state=None,
                                                         splitter='best'),
                      iid='deprecated', n_jobs=-1,
                      param_grid={'criterion': ['gini', 'entropy'],
                                   'max_depth': range(2, 32),
                                   'min_samples_leaf': range(1, 10),
                                   'min_samples_split': range(2, 10),
                                   'splitter': ['best', 'random']}},
                      pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                      scoring=None, verbose=0)
```

```
In [67]: grid_search_decision.best_params_
```

```
Out[67]: {'criterion': 'gini',
          'max_depth': 8,
          'min_samples_leaf': 8,
          'min_samples_split': 2,
          'splitter': 'best'}
```

```
In [68]: clf = DecisionTreeClassifier(criterion='entropy', max_depth=14, min_samples_leaf=2, min_samples_split=9, splitter='random')
clf.fit(train_x, train_y)
```

```
Out[68]: DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='entropy',
                                max_depth=14, max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=2, min_samples_split=9,
                                min_weight_fraction_leaf=0.0, presort='deprecated',
                                random_state=None, splitter='random')
```

```
In [69]: clf.score(test_x, test_y)
```

```
Out[69]: 0.8492872941754731
```


```
In [70]: y_pred_dt = clf.predict(test_x)
accuracy_score(test_y, y_pred_dt)
```

```
Out[70]: 0.8492872941754731
```

```
In [71]: conf_mat_dt = confusion_matrix(test_y, y_pred_dt)
true_positive_dt = conf_mat_dt[0][0]
false_positive_dt = conf_mat_dt[0][1]
false_negative_dt = conf_mat_dt[1][0]
true_negative_dt = conf_mat_dt[1][1]
Accuracy_dt = (true_positive_dt + true_negative_dt) / (true_positive_dt + false_positive_dt + false_negative_dt + true_negative_dt)
Precision_dt = true_positive_dt / (true_positive_dt + false_positive_dt)
Recall_dt = true_positive_dt / (true_positive_dt + false_negative_dt)
F1_Score_dt = 2 * (Recall_dt * Precision_dt) / (Recall_dt + Precision_dt)
auc_dt = roc_auc_score(test_y, y_pred_dt)
fpr_dt, tpr_dt, thresholds_dt = roc_curve(test_y, y_pred_dt)
print('Accuracy:: ', Accuracy_dt)
print('Precision:: ', Precision_dt)
print('Recall:: ', Recall_dt)
print('F1 Score:: ', F1_Score_dt)
print('AUC:: ', auc_dt)
```

```
Accuracy:: 0.8492872941754731
Precision:: 0.9387771520514884
Recall:: 0.8733627722475863
F1 Score:: 0.9048893024698538
AUC:: 0.7494197772737942
```

```
In [72]: plt.plot(fpr_dt, tpr_dt, color='orange', label='ROC')
plt.plot([0, 1], [0, 1], color='darkblue', linestyle='--', label='ROC curve (a
rea = %0.2f)' % auc_dt)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend()
plt.show()
```



KNN Classifier

```
In [73]: knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(train_x, train_y)
knn.score(test_x, test_y)
```

Out[73]: 0.7767879085770459

```
In [74]: y_pred_knn = knn.predict(test_x)
```

```
In [75]: print("The accuracy score is : ", accuracy_score(test_y, y_pred_knn))
```

The accuracy score is : 0.7767879085770459

```
In [76]: param_grid_knn = {'algorithm' : ['ball_tree', 'kd_tree', 'brute'],
                           'leaf_size' : [18, 20, 25, 27, 30, 32, 34],
                           'n_neighbors' : [3, 5, 7, 9, 10, 11, 12, 13]
                           }
```

```
In [103]: gridsearch_knn = GridSearchCV(knn, param_grid_knn)
```

```
In [104]: gridsearch_knn.fit(train_x, train_y)
```

```
Out[104]: GridSearchCV(cv=None, error_score=nan,  
                      estimator=KNeighborsClassifier(algorithm='ball_tree', leaf_size=  
18,  
                                                    metric='minkowski',  
                                                    metric_params=None, n_jobs=None,  
                                                    n_neighbors=13, p=2,  
                                                    weights='uniform'),  
                      iid='deprecated', n_jobs=None,  
                      param_grid={'algorithm': ['ball_tree', 'kd_tree', 'brute'],  
                                'leaf_size': [18, 20, 25, 27, 30, 32, 34],  
                                'n_neighbors': [3, 5, 7, 9, 10, 11, 12, 13]},  
                      pre_dispatch='2*n_jobs', refit=True, return_train_score=False,  
                      scoring=None, verbose=0)
```

```
In [80]: gridsearch_knn.best_params_
```

```
Out[80]: {'algorithm': 'ball_tree', 'leaf_size': 18, 'n_neighbors': 13}
```

```
In [81]: knn = KNeighborsClassifier(algorithm = 'ball_tree', leaf_size =18, n_neighbors  
=13)
```

```
In [83]: knn.fit(train_x, train_y)
```

```
Out[83]: KNeighborsClassifier(algorithm='ball_tree', leaf_size=18, metric='minkowski',  
                             metric_params=None, n_jobs=None, n_neighbors=13, p=2,  
                             weights='uniform')
```

```
In [84]: knn.score(train_x, train_y)
```

```
Out[84]: 0.8110151519808219
```

```
In [85]: knn.score(test_x, test_y)
```

```
Out[85]: 0.7970631604816908
```

```
In [86]: y_pred_knn = knn.predict(test_x)  
accuracy_score(test_y, y_pred_knn)
```


```
Out[86]: 0.7970631604816908
```



```
In [87]: conf_mat_knn = confusion_matrix(test_y, y_pred_knn)
true_positive_knn = conf_mat_knn[0][0]
false_positive_knn = conf_mat_knn[0][1]
false_negative_knn = conf_mat_knn[1][0]
true_negative_knn = conf_mat_knn[1][1]
Accuracy_knn = (true_positive_knn + true_negative_knn) / (true_positive_knn +
false_positive_knn + false_negative_knn + true_negative_knn)
Precision_knn = true_positive_knn / (true_positive_knn + false_positive_knn)
Recall_knn = true_positive_knn / (true_positive_knn + false_negative_knn)
F1_Score_knn = 2 * (Recall_knn * Precision_knn) / (Recall_knn + Precision_knn)
auc_knn = roc_auc_score(test_y, y_pred_knn)
fpr_knn, tpr_knn, thresholds_knn = roc_curve(test_y, y_pred_knn)
print('Accuracy:: ', Accuracy_knn)
print('Precision:: ', Precision_knn)
print('Recall:: ', Recall_knn)
print('F1 Score:: ', F1_Score_knn)
print('AUC:: ', auc_knn)
```

```
Accuracy:: 0.7970631604816908
Precision:: 0.9710378117457763
Recall:: 0.8039698927596083
F1 Score:: 0.8796414386182267
AUC:: 0.6029136016607196
```

```
In [88]: plt.plot(fpr_knn, tpr_knn, color='orange', label='ROC')
plt.plot([0, 1], [0, 1], color='darkblue', linestyle='--', label='ROC curve (a
rea = %0.2f)' % auc_knn)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend()
plt.show()
```



Random Forest Classifier

```
In [89]: rand_clf = RandomForestClassifier(random_state=6)
```

```
In [90]: rand_clf.fit(train_x, train_y)
```

```
Out[90]: RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                                criterion='gini', max_depth=None, max_features='auto',
                                max_leaf_nodes=None, max_samples=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=100,
                                n_jobs=None, oob_score=False, random_state=6, verbose=
                                0,
                                warm_start=False)
```

```
In [91]: rand_clf.score(test_x, test_y)
```

```
Out[91]: 0.8489800933890391
```

```
In [92]: grid_param_rdf = {
            "n_estimators" : [90,100,115],
            'criterion': ['gini', 'entropy'],
            'min_samples_leaf' : [1,2,3,4,5],
            'min_samples_split': [4,5,6,7,8],
            'max_features' : ['auto', 'log2']
        }
```

```
In [105]: grid_search_rdf = GridSearchCV(estimator=rand_clf, param_grid=grid_param_rdf,
                                           cv=5, n_jobs=-1)
```

```
In [106]: grid_search_rdf.fit(train_x, train_y)
```

```
Out[106]: GridSearchCV(cv=5, error_score=nan,
                      estimator=RandomForestClassifier(bootstrap=True, ccp_alpha=0.0,
                                                         class_weight=None,
                                                         criterion='entropy',
                                                         max_depth=None,
                                                         max_features='auto',
                                                         max_leaf_nodes=None,
                                                         max_samples=None,
                                                         min_impurity_decrease=0.0,
                                                         min_impurity_split=None,
                                                         min_samples_leaf=3,
                                                         min_samples_split=7,
                                                         min_weight_fraction_leaf=0.0,
                                                         n_estimators=100, n_jobs=None,
                                                         oob_score=False,
                                                         random_state=None, verbose=0,
                                                         warm_start=False),
                      iid='deprecated', n_jobs=-1,
                      param_grid={'criterion': ['gini', 'entropy'],
                                   'max_features': ['auto', 'log2'],
                                   'min_samples_leaf': [1, 2, 3, 4, 5],
                                   'min_samples_split': [4, 5, 6, 7, 8],
                                   'n_estimators': [90, 100, 115]},
                      pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                      scoring=None, verbose=0)
```

```
In [95]: grid_search_rdf.best_params_
```

```
Out[95]: {'criterion': 'entropy',
          'max_features': 'auto',
          'min_samples_leaf': 3,
          'min_samples_split': 7,
          'n_estimators': 100}
```

```
In [96]: rand_clf = RandomForestClassifier(criterion= 'entropy', max_features = 'auto',
                                         min_samples_leaf = 3, min_samples_split= 7,
                                         n_estimators = 100)
```

```
In [97]: rand_clf.fit(train_x, train_y)
```

```
Out[97]: RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                                criterion='entropy', max_depth=None, max_features='auto',
                                max_leaf_nodes=None, max_samples=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=3, min_samples_split=7,
                                min_weight_fraction_leaf=0.0, n_estimators=100,
                                n_jobs=None, oob_score=False, random_state=None,
                                verbose=0, warm_start=False)
```

```
In [98]: rand_clf.score(test_x, test_y)
```

```
Out[98]: 0.8585647579257802
```


```
In [99]: y_pred_rdf = rand_clf.predict(test_x)
accuracy_score(test_y, y_pred_rdf)
```

Out[99]: 0.8585647579257802

```
In [100]: conf_mat_rdf = confusion_matrix(test_y, y_pred_rdf)
true_positive_rdf = conf_mat_rdf[0][0]
false_positive_rdf = conf_mat_rdf[0][1]
false_negative_rdf = conf_mat_rdf[1][0]
true_negative_rdf = conf_mat_rdf[1][1]
Accuracy_rdf = (true_positive_rdf + true_negative_rdf) / (true_positive_rdf +
false_positive_rdf + false_negative_rdf + true_negative_rdf)
Precision_rdf = true_positive_rdf / (true_positive_rdf + false_positive_rdf)
Recall_rdf = true_positive_rdf / (true_positive_rdf + false_negative_rdf)
F1_Score_rdf = 2 * (Recall_rdf * Precision_rdf) / (Recall_rdf + Precision_rdf)
auc_rdf = roc_auc_score(test_y, y_pred_rdf)
fpr_rdf, tpr_rdf, thresholds_rdf = roc_curve(test_y, y_pred_rdf)
print('Accuracy:: ', Accuracy_rdf)
print('Precision:: ', Precision_rdf)
print('Recall:: ', Recall_rdf)
print('F1 Score:: ', F1_Score_rdf)
print('AUC:: ', auc_rdf)
```

Accuracy:: 0.8585647579257802
Precision:: 0.9483507642799678
Recall:: 0.8765615704937537
F1 Score:: 0.9110441301491614
AUC:: 0.7583667497946901

```
In [101]: plt.plot(fpr_rdf, tpr_rdf, color='orange', label='ROC')
plt.plot([0, 1], [0, 1], color='darkblue', linestyle='--', label='ROC curve (a
rea = %0.2f)' % auc_rdf)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend()
plt.show()
```



Comparing scores of different classification models

```
In [102]: print('Score of Logistic Regression::', F1_Score)
          print('Score of Decision Tree::', F1_Score_dt)
          print('Score of Random Forest::', F1_Score_rdf)
          print('Score of K-Nearest Neighbors::', F1_Score_knn)
          print('Score of XGBoost::', F1_Score_xgboost)
```

```
Score of Logistic Regression:: 0.882767554387168
Score of Decision Tree::      0.9048893024698538
Score of Random Forest::     0.9110441301491614
Score of K-Nearest Neighbors:: 0.8796414386182267
Score of XGBoost::           0.9112812342128784
```

After comparing the scores, we see that xgboost model is performing better as compared to others

Project Done By: Urvi Gadda

mailto: urvigada96@gmail.com

```
In [ ]:
```